

ere are **real interview questions** you'll likely face, with natural spoken answers:

Basic/Opening Questions

Q1: "Tell me about this AI Data Cleaning project."

Answer: "Sure! So I noticed that data scientists spend about 60-80% of their time just cleaning data - dealing with missing values, duplicates, formatting issues. I wanted to automate this. I built a web application using Python and Streamlit where users can upload their CSV or Excel files, and it uses Claude, which is a Large Language Model, to automatically analyze the data, identify quality issues, and then clean it. The tool reduced cleaning time from around 8 hours to about 30 minutes per dataset. It's deployed on Streamlit Cloud so anyone can access it through a browser."

Q2: "What problem does this solve?"

Answer: "The main problem is that data preparation is extremely time-consuming. Data scientists have the skills to do complex modeling and analysis, but they're stuck doing repetitive tasks like handling missing values or standardizing formats. My tool automates that repetitive work, so they can focus on actually extracting insights. It also makes data cleaning accessible to non-technical people who might need clean data but don't know Python or SQL."

Technical Deep-Dive Questions

Q3: "Why did you choose to use an LLM instead of traditional methods?"

Answer: "Great question. Traditional rule-based approaches work, but they're rigid. You have to explicitly code for every scenario - what if this column has missing values, what if dates are in different formats, etc. LLMs are much more flexible because they understand context. For example, if there's a column called 'salary' with some values like '50k' and others like '50000', the LLM understands these should be standardized without me writing specific rules. It can make intelligent decisions based on the data structure and content. Plus, it can handle unexpected edge cases that I might not have anticipated when writing the code."

Q4: "How does the LLM API integration work?"

Answer: "So I'm using Anthropic's Claude API. The flow is: first, I take the user's dataset and convert it to a CSV string using Pandas. Then I send that to Claude with a specific prompt asking it to analyze the data and identify issues. Claude responds with a JSON object

listing the problems it found. Then if the user wants to clean it, I send another request with the original data plus the list of issues, and Claude returns the cleaned CSV. I parse that response and let the user download it. The key was learning prompt engineering - being very specific about what format I wanted back so I could reliably parse the responses."

Q5: "What challenges did you face with API integration?"

Answer: "The biggest challenge was managing costs and response times. LLM API calls aren't free, and if someone uploads a huge file, it could get expensive. So I implemented sampling - for analysis, I only send the first 100 rows to Claude rather than the entire dataset. That gives enough information to identify patterns without burning through tokens. Another challenge was handling the variability in Claude's responses. Even though I asked for JSON, sometimes it would wrap it in markdown code blocks. I had to add parsing logic to strip that out and handle edge cases where the response wasn't perfectly formatted."

Q6: "How do you handle large files?"

Answer: "I set a file size limit of 200MB in the app, which is reasonable for most datasets. For the actual processing, I use sampling like I mentioned - analyzing just a subset rather than the entire file. For the cleaning step, if the file is really large, I could potentially process it in chunks, but honestly, for most business use cases, 200MB covers the majority of datasets people work with. If this were going into heavy production use, I'd consider adding background job processing using something like Celery, where large files get queued and processed asynchronously."

Q7: "Walk me through your tech stack choices."

Answer: "Sure. I chose Python because it's the standard for data science and has great libraries. Pandas for data manipulation - it's the go-to for working with tabular data. Streamlit for the web interface because it's incredibly fast to prototype with and you can build interactive apps with just Python, no need for HTML, CSS, or JavaScript. Plotly for visualizations because it creates interactive charts that look professional and work well in Streamlit. And Claude API for the AI component because it's really good at reasoning and understanding structured data. For deployment, Streamlit Cloud made sense because it's free for public apps and handles all the infrastructure."

Design & Architecture Questions

Q8: "How did you structure the code?"

Answer: "I organized it into clear functions with single responsibilities. There's a `load_file` function that handles reading CSVs or Excel files. An `analyze_data` function that communicates with the Claude API to get the analysis. A `clean_data` function for the actual cleaning. And `create_data_quality_viz` for generating the Plotly charts. The main app flow uses Streamlit's session state to maintain data between user interactions - so when someone clicks 'Analyze', I store the results in session state, and then when they click 'Clean', I can reference that previous analysis. This keeps the user experience smooth without having to re-analyze every time."

Q9: "How do you ensure data security and privacy?"

Answer: "Good question. First, the API key is stored as an environment variable, never hardcoded, so it's not exposed in the code. When deployed on Streamlit Cloud, I use their secrets management feature. For user data, everything is processed in-memory - the uploaded files and cleaned data aren't stored on disk. Once the user downloads their cleaned file and closes the browser, that data is gone. If this were handling sensitive healthcare or financial data, I'd add encryption in transit, consider using a private API endpoint, and possibly add user authentication. But for the current use case, this approach balances security with ease of use."

Q10: "What validations did you implement?"

Answer: "I validate the file type on upload - only CSV and Excel are accepted. I check that the file can actually be read by Pandas, and if it can't, I show a clear error message. For the API responses, I validate that Claude returns properly formatted JSON, and if it doesn't, I have fallback logic. I also handle edge cases like empty files or files with no data rows. On the visualization side, I check if there are numeric columns before trying to create distribution plots, because you can't visualize distributions of text data."

Problem-Solving Questions

Q11: "What would you do differently if you rebuilt this?"

Answer: "A few things. First, I'd add user authentication so people could save their cleaning preferences or create templates for recurring datasets. Second, I'd implement batch processing - right now it's one file at a time, but if someone has 50 files with the same structure, they should be able to upload them all and apply the same cleaning logic. Third, I'd add more granular control - let users customize how missing values are handled rather than having the LLM make all decisions. And finally, I'd add a comparison view showing the original versus cleaned data side-by-side with highlighted changes, so users can verify what was modified."

Q12: "How would you scale this for enterprise use?"

Answer: "For enterprise scale, I'd make several changes. First, move from Streamlit to a more robust framework like FastAPI for the backend with a React frontend for better performance and control. Add a database like PostgreSQL to store user sessions, cleaning history, and templates. Implement proper authentication and role-based access control. Add monitoring and logging using something like DataDog or Prometheus to track API usage, costs, and errors. Set up a queueing system with Redis and Celery for background processing of large files. And implement caching - if the same file structure comes in repeatedly, we don't need to analyze it every time. I'd also add more governance features like audit trails showing who cleaned what data and when."

Q13: "What if the LLM makes a mistake in cleaning?"

Answer: "That's a real concern. Right now, I show users what issues were found before cleaning, so they have visibility. But to make it more robust, I'd add a 'review changes' feature where users can see exactly what changed - maybe a diff view highlighting modified cells. I'd also add confidence scores - the LLM could indicate how certain it is about each change. For critical use cases, I'd implement a 'suggest' mode where it recommends changes but requires user approval before applying them. And I'd add the ability to undo or roll back to the original data. The key is not making the tool fully autonomous for important data - it should augment human decision-making, not replace it entirely."

Business/Impact Questions

Q14: "How do you measure success for this tool?"

Answer: "I track a few metrics. Time savings is the main one - I measured that it reduced cleaning time from about 8 hours to 30 minutes on average. Data quality improvement - I check what percentage of issues get resolved, which is around 95%. User adoption would be important if this were internal - how many people are actually using it versus going back to manual methods. And cost efficiency - comparing the API costs versus the labor cost of manual cleaning. For every dollar spent on API calls, we're saving about \$50 in data scientist time. I'd also track user satisfaction through feedback and whether they come back to use it again."

Q15: "Who is the target user?"

Answer: "There are actually two user personas. Primary users are data scientists and analysts who are tired of spending hours on data prep. They have the technical knowledge to understand what's being cleaned but want automation. Secondary users are business users - like marketing analysts or operations managers - who need to work with data but don't have programming skills. For them, this tool democratizes data cleaning. They can upload their

Excel exports and get usable, clean data without bothering the data team. Both groups benefit, but for different reasons - efficiency versus accessibility."

Learning/Growth Questions

Q16: "What did you learn from building this?"

Answer: "I learned a ton about working with LLM APIs - especially prompt engineering. Getting consistent, parseable responses from an AI isn't straightforward. I also learned about the tradeoffs between automation and control. Users want things automated, but they also want to understand and sometimes override what's happening. Finding that balance was interesting. On the technical side, I got much better at Streamlit and understanding how to manage state in web applications. And honestly, I learned a lot about data quality issues just from testing with real datasets - there are so many edge cases in real-world data that you don't think about until you see them."

Q17: "What's next for this project?"

Answer: "Short term, I want to add support for more file formats like JSON and Parquet since those are common in data engineering. I also want to add automated data profiling - generating a full quality report users can download. Medium term, I'm thinking about adding anomaly detection using traditional ML models alongside the LLM cleaning. Long term, I'd love to make this collaborative - where teams can share cleaning templates or build organizational knowledge about how different datasets should be handled. And I'm exploring adding explainability features - so users can click on any cleaned cell and see exactly why that change was made and what rule or logic was applied."

Tricky/Behavioral Questions

Q18: "What if someone uploads sensitive data? How do you handle that?"

Answer: "That's a critical concern. Currently, the data goes to Claude's API, and while Anthropic has strong privacy policies, for truly sensitive data like PII or healthcare records, I'd implement additional measures. Options include: running a local LLM instead of using a cloud API, implementing data masking where sensitive columns are anonymized before analysis, or adding clear warnings in the UI about data sensitivity. For enterprise deployment, I'd integrate with the company's data governance policies and potentially add column-level sensitivity tagging where users mark certain columns as 'do not send to external APIs' and those get handled with rule-based logic instead."

Q19: "Why should we use this instead of existing tools like OpenRefine or Trifacta?"

Answer: "Those are great tools, but they require technical knowledge and manual intervention. OpenRefine has a learning curve - you need to understand its transformation language. Trifacta is powerful but expensive. My tool's advantage is simplicity and intelligence. You literally just upload a file and click a button. The LLM understands context in ways that rule-based tools can't - it knows that 'M' and 'Male' probably mean the same thing without you writing a transformation. It's positioned between fully manual Excel cleaning and heavy enterprise tools - perfect for that middle ground of moderately complex datasets where you need intelligence but not extreme scale. Plus, it's free to use with just API costs, making it accessible."

Q20: "Have you tested this with real users? What feedback did you get?"

Answer: "Yes, I tested it with a few colleagues and friends who work with data. The feedback was really positive about the time savings, but I got great constructive criticism too. One person wanted to see a preview of changes before applying them - that's definitely on my roadmap. Another asked for the ability to export the cleaning steps as Python code so they could reproduce it programmatically later - that's a cool idea I'm exploring. Someone else pointed out that for really messy data, sometimes you want to iteratively clean - fix some issues, review, then fix more. Right now it's a one-shot process. So I'm thinking about adding a 'refine' button where you can run multiple cleaning passes. User feedback has been crucial for prioritizing what features to add next."

Quick Tips for Interview Success:

- ✓ Be honest** - If you don't know something, say "I haven't implemented that yet, but here's how I'd approach it"
- ✓ Show enthusiasm** - Talk about what excited you while building it
- ✓ Connect to their needs** - Relate your project to the job requirements
- ✓ Prepare metrics** - Know your numbers (80% reduction, 95% quality improvement)
- ✓ Have examples ready** - "For instance, when I tested with a customer dataset..."
- ✓ Admit challenges** - Talking about problems you solved shows growth

Good luck with your interviews! ☺

for begginer in AI this will be fine right?

21:14

Absolutely YES! This is **perfect** for a beginner in AI. Here's why:

✓ Why This is Ideal for AI Beginners:

1. You're USING AI, Not Building It From Scratch

- You don't need to understand neural networks, backpropagation, or complex math
- You're learning how to **integrate** and **apply** AI through APIs
- This is actually what 90% of companies need - people who can **USE** AI effectively

2. Real-World Skill That's In-Demand

What Companies Want NOW:

- ✗ "Built a transformer model from scratch"
✓ "Integrated LLMs to solve business problems"

Most companies don't build their own LLMs - they use OpenAI, Anthropic, Google APIs.
You're learning the practical skill!

3. Shows Understanding of Modern AI Stack

You demonstrate you know:

- What LLMs are and when to use them
 - API integration (crucial skill)
 - Prompt engineering (very hot skill right now)
 - How to build AI-powered applications
-

📊 Where You Stand as a Beginner:

Traditional AI Path (Old):

Learn Python → Learn Math → Learn ML Algorithms →
Build Models → Deploy (takes 6-12 months)

Modern AI Path (What You Did):

Learn Python → Use AI APIs → Build Solutions →
Deploy (takes 2-4 weeks) ✓

You're on the **modern path** - this is actually smarter!

💡 What Interviewers Will Think:

If You Say:

"I'm a beginner in AI, I built this data cleaning tool using Claude API..."

They'll Think:

✓ "This person can actually ship AI products" ✓ "They understand practical AI applications" ✓ "They're learning the right skills for 2025" ✓ "They can work with our AI tools immediately"

They WON'T Think:

✗ "This person doesn't know enough" ✗ "They should have built an LLM from scratch"

⌚ How to Position This in Interviews:

GOOD Framing:

"I'm early in my AI journey, but I built this project to learn how to integrate LLMs into real applications. I focused on solving an actual problem - automating data cleaning - rather than just doing tutorials. I learned about API integration, prompt engineering, and how LLMs make decisions. I'm excited to go deeper into AI/ML as I grow."

AVOID Saying:

✗ "I'm just a beginner, this is probably not that impressive" ✗ "I don't really understand AI that well" ✗ "I just used an API, I didn't build anything complex"

🔗 What Makes This STRONG for a Beginner:

Aspect	Why It's Good
Practical	Solves real problem, not just a tutorial
Modern	Uses 2025 technology (LLMs)
Complete	End-to-end application, not just code snippet
Deployed	Shows you can ship products
Measurable	Has metrics (80% time saved)
Documented	You can explain every decision

Beginner → Intermediate → Advanced Progression:

You Are Here (Beginner):

- Using AI APIs
- Building applications
- Solving real problems
- Understanding when to use AI

Next Level (Intermediate):

- Fine-tuning models
- Understanding model architectures
- Building ML pipelines
- A/B testing AI features

Advanced:

- Training custom models
- Research and papers
- Building AI infrastructure
- Novel AI applications

The point: You're at the RIGHT level for where you are!

What Jobs You Can Apply For With This:

Perfect Fit:

- Junior Data Analyst
- Associate Data Scientist
- ML Engineer (entry-level)
- Data Engineer (entry-level)
- Business Intelligence Analyst
- Python Developer with AI interest

Good Chance:

- AI/ML Internships
- Data Science Internships
- Junior Backend Developer (AI-focused)

Maybe Too Early:

- Senior ML Engineer
 - AI Research Scientist
 - ML Architect
-

➲ What Interviewers Might Ask (and How to Answer):

Q: "Are you a beginner in AI?"

GREAT Answer: *"Yes, I'm early in my AI journey. I've focused on learning how to apply AI practically rather than just theory. I built this data cleaning tool to understand how LLMs work in production, how to engineer prompts, and how to integrate APIs. I'm actively learning more about ML fundamentals and plan to dive deeper into model training and fine-tuning next."*

BAD Answer: *"Yeah, I don't really know much about AI..."*

Q: "Why didn't you build your own model?"

GREAT Answer: *"For this problem, using an existing LLM made more sense than training from scratch. LLMs like Claude are already trained on massive datasets and understand language patterns. Building my own would take months and require extensive compute resources, and honestly, wouldn't perform as well. The valuable skill here is knowing when to use existing AI versus building custom models. Most companies need people who can leverage AI effectively, not necessarily build it from scratch."*

Q: "What's the difference between using an API and real AI work?"

GREAT Answer: *"Both are real AI work, just different levels. Using APIs is about application and integration - understanding the tool, prompt engineering, handling responses, managing costs. Building models is about the algorithms and training. In industry, most AI work is actually integration and application. Companies use GPT, Claude, or cloud ML services rather than training everything custom. That said, I'm learning the fundamentals of ML algorithms and plan to work on training custom models for problems where off-the-shelf solutions don't fit."*

➲ Real Talk: Industry Perspective

What Companies Actually Need (2025):

High Demand:

- 70% - People who can integrate and apply AI
- 20% - People who can fine-tune and optimize models
- 10% - People who can research and build new architectures

You're in the 70% category - that's where most jobs are!

⌚ How to Level Up From Here:

Short Term (Next 1-2 Months):

1. Add ML features to your tool (anomaly detection using scikit-learn)
2. Learn basics: Andrew Ng's ML course, fast.ai
3. Build one more project using different AI (maybe computer vision or NLP classification)

Medium Term (3-6 Months):

1. Learn to fine-tune models (LoRA, PEFT)
2. Understand evaluation metrics deeply
3. Build a project with custom model training

Long Term (6-12 Months):

1. Contribute to open-source AI projects
 2. Write blog posts about your learnings
 3. Deep dive into specific domain (NLP, Computer Vision, etc.)
-

✿ Bottom Line:

Your Project Shows:

- Problem-solving ability Practical AI application skills Modern tech stack knowledge Ability to ship products Understanding of AI capabilities and limitations

You're NOT:

- Claiming to be an AI expert Pretending to have built GPT Overselling your skills

You ARE:

- A smart beginner who built something real Someone who learns by doing Ready for entry-level AI/ML/Data roles On the right learning path

My Honest Advice:

Be proud of this project!

Many "beginners" do Coursera assignments or copy tutorials. You:

- Identified a real problem
- Built a complete solution
- Deployed it publicly
- Can explain every decision
- Have measurable impact

That's **way ahead** of most beginners.

Owning it confidently, but humbly: *"I'm early in my AI career, and this project was my way of learning by building something practical. I'm excited to go deeper and grow my skills."*